





# Projet Analyse De Donnée

L3

**Hossein Khani**



# Content:

- Why Data Analyses
- Data Manipulation (Pandas Library)
- Data Visualisation (Matplotlib, Pyplot, Seaborn)
- Linear Regression
- Principle Component Analysis
- Non-Negative Matrix Factorization
- Orthogonal Matching pursuit

# NEEDS:

- Basic Python Skills (Lists, Dictionaries, Functions, methods,....)
- Working with DataFrames (Data Cleaning and manipulation with Pandas Library)
- Working with Matrices (Numpy Library)
- Mathematics behind Machine learning Techniques (Mostly probability and statistics)
- Machine learning library (Scipy or Sklearn)

# QUESTION:

Input:

longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_proximity
-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	452600.0	NEAR BAY
-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0	NEAR BAY
-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0	NEAR BAY
-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0	NEAR BAY
-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0	NEAR BAY


Output:

-121.32	39.43	18.0	1860.0	409.0	741.0	349.0	1.8672	??	INLAND
---------	-------	------	--------	-------	-------	-------	--------	----	--------

REGRESSION

# QUESTION:

Input:



species	margin1	margin2	margin3	margin4	margin5	margin6	margin7	margin8	...	texture55	texture56	texture57	texture58
Acer_Opalus	0.007812	0.023438	0.023438	0.003906	0.011719	0.009766	0.027344	0.0	...	0.007812	0.000000	0.002930	0.002930
Pterocarya_Stenoptera	0.005859	0.000000	0.031250	0.015625	0.025391	0.001953	0.019531	0.0	...	0.000977	0.000000	0.000000	0.000977
Quercus_Hartwissiana	0.005859	0.009766	0.019531	0.007812	0.003906	0.005859	0.068359	0.0	...	0.154300	0.000000	0.005859	0.000977

Output:

??

0.000000 0.003906 0.023438 0.005859 0.021484 0.019531 0.023438 0.0 ... 0.000000 0.000977 0.000000 0.000000

CLASSIFICATION

# Install Python

- <https://www.python.org/>
  - pip Python package management system : **python3 -m pip --version**
  - install jupyter notebook: **python3 -m pip install -U jupyter**
  - Install pandas: **pip install pandas**
- The Jupyter Notebook is the original web application for creating and sharing computational documents.
  - Pandas the main tool of data analyse
  - Pandas permits us to import data from various sources for example (CSV), and manipulate them.

# DataFrames:

➤ <https://insights.stackoverflow.com/survey>

How to use Pandas to work with DataFrame.....

1. How to read data from csv file,
2. Take a look at the dataframe,
3. Where dataframe comes from, its equivalent in python
4. Series objects and accessing multi-columns
5. Indexing
6. Accessing rows in DataFrames
7. Setting index for data frame
8. Changing columns' names
9. Changing single row's values



# Numpy:

As a Data Analyst how to collect data?

## ➤ List?

- Collection of values
- Hold different types
- Change, add, remove

## ➤ What we need more?

- Mathematical operations over collections
- Speed

# Numpy:

Body mass Index:

```
Height = [1.73, 1.68, 1.71, 1.89, 1.79]  
Weight = [65.4, 59.2, 63.6, 88.4, 68.7]
```

```
Weight / Height ** 2
```

```
-----  
TypeError                                Traceback (most recent call last)  
<ipython-input-43-0f6f8ba4f85f> in <module>  
----> 1 Weight / Height ** 2  
  
TypeError: unsupported operand type(s) for ** or pow(): 'list' and 'int'
```

**To Solve:** Looping over elements?

→ Not fast and efficient

# Numpy (numeric python):

## Solution?

numpy arrays:

- Alternative to python lists
- Calculations over entire arrays
- Easy and Fast

To Install:

**pip3 install numpy**

# Numpy

```
import numpy as np
```

```
np_height = np.array(Height)  
np_height
```

```
array([1.73, 1.68, 1.71, 1.89, 1.79])
```

```
np_weight = np.array(Weight)  
np_weight
```

```
array([65.4, 59.2, 63.6, 88.4, 68.7])
```

```
bmi = np_weight / np_height**2  
bmi
```

```
array([21.85171573, 20.97505669, 21.75028214, 24.7473475 , 21.44127836])
```

Python is able to treat numpy arrays as single elements.

Where the speed comes from?

Numpy arrays collect values of the same type:

- Either integer
- Either float
- String
- .....

# Numpy (Remarks)

```
np.array([1.0 , "Hossein" , True])
```

```
array(['1.0', 'Hossein', 'True'], dtype='<U32')
```

- Numpy array, is a data type in python.
- It has its own methods.
- These methods might act differently on arrays compared to other types.

# Numpy (Remarks)

Example:

```
python_list = [1,2,3]  
numpy_array = np.array([1,2,3])
```

```
python_list + python_list
```

```
[1, 2, 3, 1, 2, 3]
```

```
numpy_array+numpy_array
```

```
array([2, 4, 6])
```

# Numpy (Subsetting)

## Example:

```
bmi
```

```
array([21.85171573, 20.97505669, 21.75028214, 24.7473475 , 21.44127836])
```

```
bmi[2]
```

Referring to specific index

```
21.750282138093777
```

```
bmi > 21
```

Looking for specific values

```
array([ True, False,  True,  True,  True])
```

```
bmi[bmi<21]
```

```
array([20.97505669])
```



# Numpy (2D)

```
type(np_height)
```

```
numpy.ndarray
```

```
np_2d = np.array([[1.73, 1.68, 1.71, 1.89, 1.79],  
                  [65.4, 59.2, 63.6, 88.4, 68.7]])
```

```
np_2d
```

```
array([[ 1.73,  1.68,  1.71,  1.89,  1.79],  
       [65.4, 59.2, 63.6, 88.4, 68.7]])
```

```
np_2d.shape
```

```
(2, 5)
```

# Numpy (2D)

```
np_2d = np.array([[1.73, 1.68, 1.71, 1.89, 1.79],  
                  [65.4, 59.2, 63.6, 88.4, 68.7]])
```

```
np_2d
```

```
array([[ 1.73,  1.68,  1.71,  1.89,  1.79],  
       [65.4,  59.2,  63.6,  88.4,  68.7]])
```

```
np_2d[0]
```

```
array([1.73, 1.68, 1.71, 1.89, 1.79])
```

```
np_2d[0][2]
```

```
1.71
```

```
np_2d[0,2]
```

```
1.71
```

```
np_2d[0,1:3]
```

```
array([1.68, 1.71])
```

Two ways to select



# Numpy (Basic Statistics)

```
np_2d = np.array([Height,  
                  Weight])
```

```
np_2d
```

```
array([[ 1.73,  1.68,  1.71,  1.89,  1.79],  
       [65.4 , 59.2 , 63.6 , 88.4 , 68.7 ]])
```

```
np.mean(np_2d[0,:])
```

```
1.7600000000000002
```

```
np.median(np_2d[0,:])
```

```
1.73
```

```
np.sum(np_2d[0,:])
```

```
8.8
```

# Numpy (Data Generation)

```
height = np.round(np.random.normal(1.75, 2.0, 5000), 2)  
weight = np.round(np.random.normal(10.32, 15.0, 5000), 2)
```

```
np_city = np.column_stack((height, weight))
```

```
np_city.shape
```

```
(5000, 2)
```

# Numpy (Data Generation)

```
np.zeros([4,5],dtype = int)
```

```
array([[0, 0, 0, 0, 0],  
       [0, 0, 0, 0, 0],  
       [0, 0, 0, 0, 0],  
       [0, 0, 0, 0, 0]])
```

```
np.ones([4,5], dtype = int)
```

```
array([[1, 1, 1, 1, 1],  
       [1, 1, 1, 1, 1],  
       [1, 1, 1, 1, 1],  
       [1, 1, 1, 1, 1]])
```

```
np.full((2,3), 6, dtype = int)
```

```
array([[6, 6, 6],  
       [6, 6, 6]])
```

# Numpy (Dtype)

- Python types: int, float, bool,...  
Their size depends on the platform they are applied to...
- Dtypes: numpy numerical types are instances of **dtype** objects. The numpy types have fixed-sizes.

np.int32, np.int64, np.bool8, np.float32, np.float64

```
z = np.zeros([2,3], dtype = np.bool8)
```

```
z
```

```
array([[False, False, False],  
       [False, False, False]])
```


```
type(z)
```

```
numpy.ndarray
```

```
z.dtype
```

```
dtype('bool')
```

# Data Visualisation:




# The most important visualization library :

## Matplotlib:

```
import matplotlib.pyplot as plt
```


```
years = [1950, 1970, 1990, 2010]  
pop = [2.519, 3.692, 5.263, 6.972]
```

```
plt.plot(years , pop)  
plt.show()
```





plt.plot() :




```
year = [1950 , 1970 , 1990 , 2010]  
pop  = [2.519, 3.692, 5.263, 6.972]
```

# plt.scatter() :

```
years = [1950, 1970, 1990, 2010]  
pop = [2.519, 3.692, 5.263, 6.972]
```

```
plt.scatter(years , pop)  
plt.show()
```



Scatter plot is used when we need to measure the **correlation** between two attributes.

plt.scatter() :

```
np.corrcoef(years, pop)
```

```
array([[1.          , 0.99664316],  
       [0.99664316, 1.          ]])
```


```
import scipy.stats as st  
st.pearsonr(years, pop)
```

```
(0.996643163032238, 0.0033568369677620113)
```

Correlation


P-value

# P-Value:



Probabilities?

# P-Value:



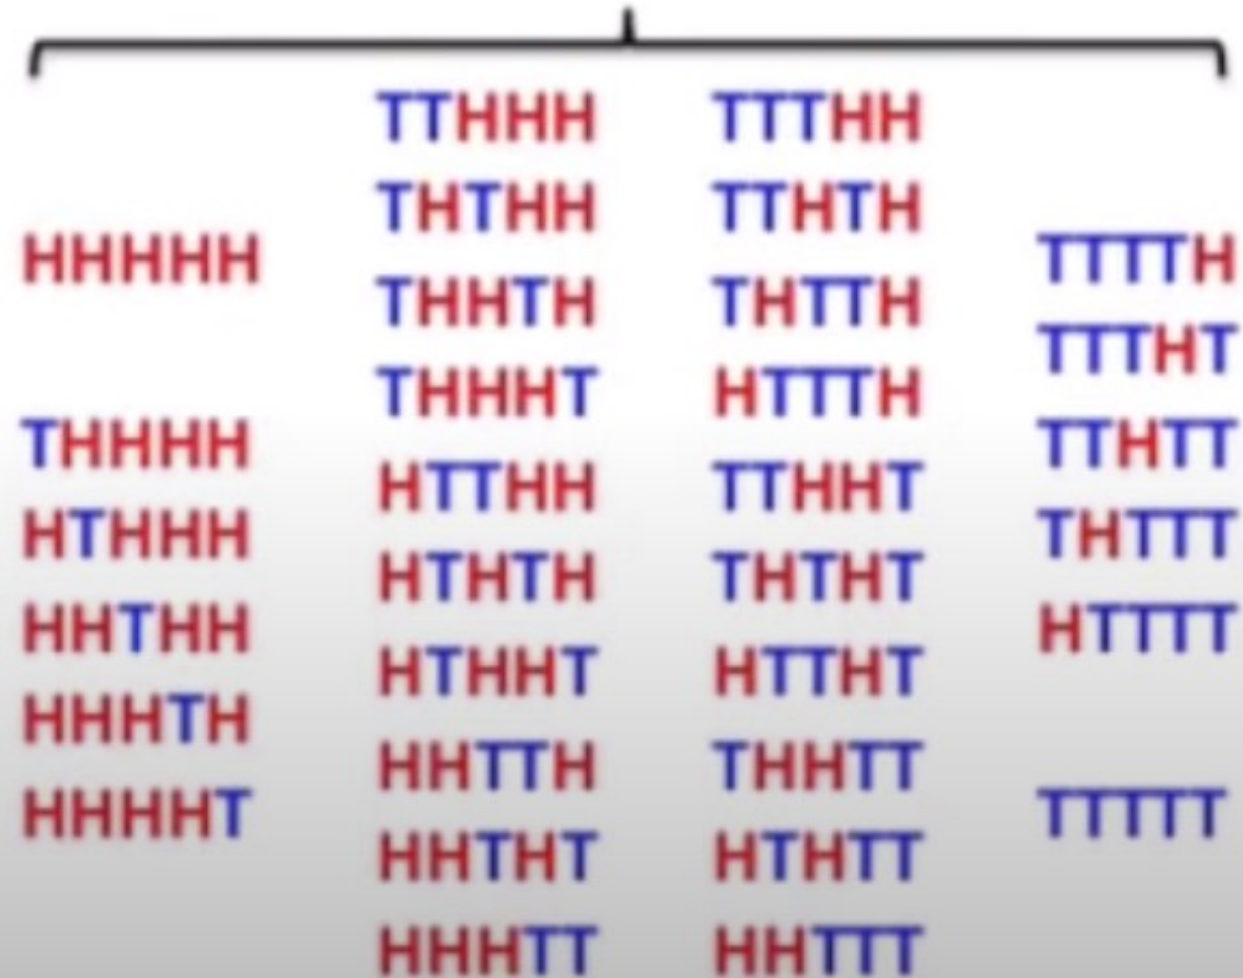
# P-Value (flipping a coin 5 times):

$$p(HHHHH) = \frac{1}{32}$$

$$p(TTTTT) = \frac{1}{32}$$

$$p\text{-value}(HHHHH) = \frac{1}{32} + \frac{1}{32} = 0.0625$$

Outcomes



	TTTTHH	TTTTHH	
	THTHH	TTHTH	
HHHHH	THHTH	THTTH	TTTTH
	THHHT	HTTTH	TTTHT
THHHH	HTTTH	TTHHT	TTHTT
HTHHH	HTHTH	THTHT	THTTT
HHTHH	HTHHT	HTTHT	HTTTT
HHHHT	HHTTH	THHTT	
	HHTHT	HTHTT	TTTTT
	HHHTT	HHTTT	


**P-Value:** probability that random chance generated the data or something else that is equal or rarer.

**P-values:**

$$\frac{1}{4} + \frac{1}{4} = \frac{1}{2}$$

1


$\frac{1}{2}$



# plt.scatter() :

```
years = [1950, 1970, 1990, 2010]  
pop = [2.519, 3.692, 5.263, 6.972]
```

```
plt.scatter(years , pop)  
plt.show()
```



Scatter plot is used when we need to measure the **correlation** between two attributes.




# plt.hist() :



```
values = [0,0.6,1.4,1.6,2.2,2.5,2.6,3.2,3.5,3.9,4.2,6]
```

```
plt.hist(values , bins = 3)
```

```
(array([4., 6., 2.]),  
 array([0., 2., 4., 6.]),  
 <BarContainer object of 3 artists>)
```



plt.hist() :



# Distribution of Data:

- Which is the most frequent data? `statistics.mode()`
- The data is centered around which point? `Numpy.mean()`
- What is the value observed in 50% of the time? `Numpy.median()`
- How vary the values are ? `np.std()`

# Distribution of Data:

Most of the time it takes 80 mins  
Half of the times it takes 80 mins  
On average it takes 80 mins



How long does it take  
to go from  
City A to city B



# Distribution of Data:

```
values
```

```
[0, 0.6, 1.4, 1.6, 2.2, 2.5, 2.6, 3.2, 3.5, 3.9, 4.2, 6]
```

```
import numpy as np  
np.mean(values)
```


```
2.6416666666666666
```

```
np.median(values)
```


```
2.55
```

```
import statistics as sts  
sts.mode(values)
```


```
0
```



mode




median




mean

# Distribution of Data:




Left-Skewed (Negative Skewness)



Right-Skewed (Positive Skewness)

# Distribution of Data:

```
plt.hist(df['Fare'], bins = 30)  
plt.show()
```




$$z = \frac{x - \mu}{\sigma}$$

$\mu$  = Mean

$\sigma$  = Standard Deviation

```
plt.hist(np.std(df['Fare']), bins = 30)  
plt.show()
```



# Customization:


- Add labels to the axis: `plt.xlabel()` , `plt.ylabel()`
- Add Title to the plot : `plt.title()`
- Changing values on the axis: `plt.xticks()` , `plt.yticks()`
- Labeling values on the axis



# Customization:

```
years = [1950, 1970, 1990, 2010]
pop = [2.519, 3.692, 5.263, 6.972]
np_years = np.array(years)
np_pop = np.array(pop)
plt.plot(np_years, np_pop)
plt.xlabel('year')
plt.ylabel('population')
plt.title('World population in years')
```


```
Text(0.5, 1.0, 'World population in years')
```



# Customization:

```
years = [1950, 1970, 1990, 2010]
pop = [2.519, 3.692, 5.263, 6.972]
np_years = np.array(years)
np_pop = np.array(pop)
plt.plot(np_years, np_pop)
plt.xlabel('year')
plt.ylabel('populattion')
plt.yticks([0, 2, 4, 6, 8])
plt.title('World population in years')
```


```
Text(0.5, 1.0, 'World population in years')
```



# Customization:

```
years = [1950, 1970, 1990, 2010]
pop = [2.519, 3.692, 5.263, 6.972]
np_years = np.array(years)
np_pop = np.array(pop)
plt.plot(np_years, np_pop)
plt.xlabel('year')
plt.ylabel('populattion')
plt.yticks([0, 2, 4, 6, 8])
plt.title('World population in years')
```


```
Text(0.5, 1.0, 'World population in years')
```



# Customization:

```
years = [1950, 1970, 1990, 2010]
pop = [2.519, 3.692, 5.263, 6.972]
np_years = np.array(years)
np_pop = np.array(pop)
plt.plot(np_years, np_pop)
plt.xlabel('year')
plt.ylabel('populattion')
plt.yticks([0, 2, 4, 6, 8], ['B', '2B', '4B', '6B', '8B'])
plt.title('World population in years')
```

Text(0.5, 1.0, 'World population in years')



# Machine Learning:


How you write a code with traditional programming technique to detect spams?

- What a spam looks like, what are the patterns,
- Write a detection algorithm for each pattern,... .

Problem??

There is an infinite number of patterns!


# ML (Supervised):



In ML, the model will learn (based on some examples) which patterns are representative of a spam.

Classification

# ML (Supervised):



Given a set of Instances and their corresponding value, we can guess what is the value of a newly entered instance.

Regression

# ML (Example):


	Country	GDP per capita	Life satisfaction
0	Australia	50961.865	82.1
1	Austria	43724.031	81.0
2	Belgium	40106.632	80.5
3	Brazil	8669.998	73.7
4	Canada	43331.961	81.5
5	Chile	13340.905	78.9
6	Czech Republic	17256.918	78.2

Given a GDP per capita in a country, can you guess what is the life satisfaction index?



# ML (Example):

```
import numpy as np
plt.scatter(data['GDP per capita'], data['Life satisfaction'])
x = np.array([1000, 100000])
plt.xlabel('GDP per capita')
plt.ylabel('Life satisfaction')
plt.show()
```



What is the simplest and common pattern in the scatter plot?


$$y = \theta_0 + \theta_1 x$$

# ML (Example):

$\theta_0$

$\theta_1$

```
import numpy as np
plt.scatter(data['GDP per capita'] , data['Life satisfaction'])
x = np.array([1000 , 100000])
t_0 = 73
t_1 = 8.9e-05
plt.plot(x , t_0 + t_1*x , c = 'red')
plt.xlabel('GDP per capita')
plt.ylabel('Life satisfaction')
plt.show()
```




*life satisfaction*  
 $= \theta_0 + \theta_1 * GDP\ per\ Capita$

# ML (Example):

$\theta_0$  →  
 $\theta_1$  →

```
import numpy as np
plt.scatter(data['GDP per capita'] , data['Life satisfaction'])
x = np.array([1000 , 100000])
t_0 = 76.61443338
t_1 = 8.82017196e-05
plt.plot(x , t_0 + t_1*x , c = 'red')
plt.xlabel('GDP per capita')
plt.ylabel('Life satisfaction')
plt.show()
```



$$\begin{aligned} \text{life satisfaction} \\ = \theta_0 + \theta_1 * \text{GDP per Capita} \end{aligned}$$

# ML (Linear Assumption):

	GDP per capita	Life satisfaction	
$x^{(1)}$	50961.865	82.1	$y^{(1)}$
$x^{(2)}$	43724.031	81.0	$y^{(2)}$
$x^{(3)}$	40106.632	80.5	$y^{(3)}$
$x^{(4)}$	8669.998	73.7	$y^{(4)}$
	43331.961	81.5	
	13340.905	78.9	
	17256.918	78.2	
	52114.165	80.1	
	17288.083	76.5	
	41973.988	80.7	

$$\hat{y}^{(i)} = \theta_0 + \theta_1 \times x^{(i)}$$

# ML (Example):

**Main assumption:** The data follows a linear model:

$$\textit{life satisfaction} = \theta_0 + \theta_1 * \textit{GDP per Capita}$$

➤ How you know which values make your model perform best?


- Fitness Function
- Cost Function (typically used for linear regression problems.)

**Linear Regression algorithm comes into play:**

you feed it your training examples and it finds the parameters that make the linear model fit best to your data.  
This is called *training* the model.


# ML (Example):

Cost Function:



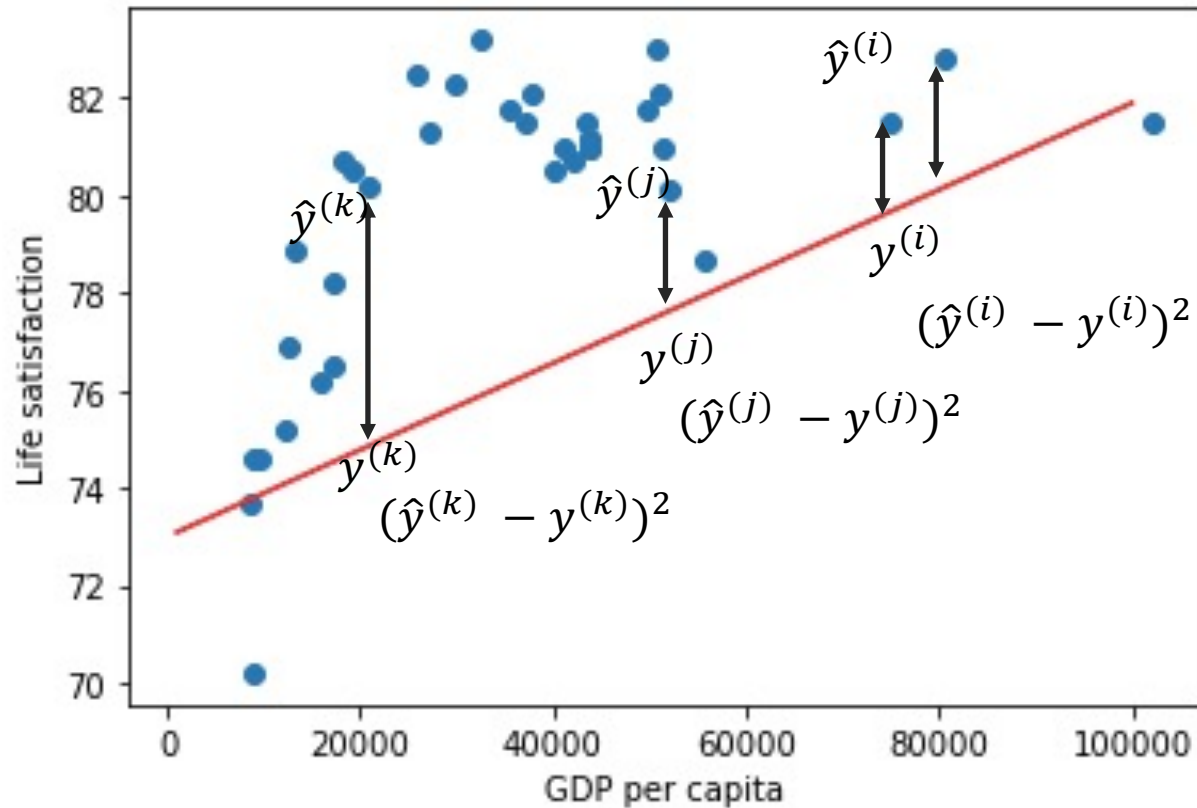
# ML (Example):

Cost Function:



# ML (Example):

Cost Function: Root Mean Square Error



$$\text{RMSE} = \sqrt{\frac{1}{m} \sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)})^2}$$

Always positive



# ML (Example):

Sklearn (Python library for sklearn)

```
from sklearn.linear_model import LinearRegression
```

```
model = LinearRegression()
```

```
model.fit(data[['GDP per capita']], data[['Life satisfaction']])
```

Model Training



# ML (Example):

Sklearn (Python library for sklearn)

```
model.predict(data[['GDP per capita']])
```

```
array([[81.10935751],  
       [80.47096811],  
       [80.15190729],  
       [77.37914212],  
       [80.43638686],  
       [77.79112415],  
       [78.13652323],  
       [81.21099235],  
       [78.13927203],  
       [80.3166113 ],  
       [79.9374337 ],  
       [80.23039615],  
       [78.20773465],  
       [77.69401308],  
       [81.09989505],  
       ...])
```

Model Prediction

```
data[['Life satisfaction']]
```

Life satisfaction	
0	82.1
1	81.0
2	80.5
3	73.7
4	81.5
5	78.9
6	78.2
7	80.1
8	76.5
9	80.7
10	82.1
11	81.0
12	80.7
13	75.2
14	83.0
15	81.0

# ML (Example):

Predictor (A combination of attributes)

GDP per capita Life satisfaction

Target variable

Samples

50961.865	82.1
43724.031	81.0
40106.632	80.5
8669.998	73.7
43331.961	81.5
13340.905	78.9
17256.918	78.2
52114.165	80.1
17288.083	76.5
41973.988	80.7

Training (Minimizing RMSE)

Predictions

```
[[[81.10935751],  
 [80.47096811],  
 [80.15190729],  
 [77.37914212],  
 [80.43638686],  
 [77.79112415],  
 [78.13652323],  
 [81.21099235],  
 [78.13927203],  
 [80.3166113 ]],
```

Assumption of Linearity

Note: In general there might be more than one attribute:

- In this case, the first attribute of sample (i) is represented by variable  $x_1^{(i)}$ ,
- The second attribute would be  $x_2^{(i)}$
- ....
- The attribute p would be  $x_p^{(i)}$

The linear assumption:  $\hat{y}^{(i)} = \theta_0 + \theta_1 \times x_1^{(i)} + \theta_2 \times x_2^{(i)} + \dots + \theta_p \times x_p^{(i)}$

Hyperplane

Note: in some texts instead of  $\hat{y}^{(i)}$ , they use  $h_{\theta}(x^i)$

# ML (Example):

How the training part works? (The minimization of RMSE)

$$\text{Min} \quad \text{RMSE} = \sqrt{\frac{1}{m} \sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)})^2}$$

Is equal to  $\text{Min} \quad (\theta^T X - y)^2$

$$\theta = [\theta_0 \quad \theta_1 \quad \theta_2 \quad \dots \quad \theta_p]$$

$$X = \begin{pmatrix} 1 & x_1^1 & x_2^1 & \dots & x_p^1 \\ 1 & x_1^2 & x_2^2 & \dots & x_p^2 \\ \vdots & \vdots & \dots & \vdots & \\ 1 & x_1^n & x_2^n & \dots & x_p^n \end{pmatrix}$$

$$y = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ y^{(3)} \\ \vdots \\ y^{(n)} \end{bmatrix}$$

# ML (Example):

How the training part works? (The minimization of RMSE)

$$\text{Min} \quad (\theta^T X - y)^2$$

$$\arg \min_{\theta \in \mathbb{R}^{p+1}} (\theta^T X - Y)^T (\theta^T X - Y)$$

$$\nabla_{\theta} (\theta^T X - Y)^T (\theta^T X - Y) = 0$$

$$-2 X^T (y - \theta^T X) = 0$$

**Normal Equation**

$$\theta = (X^T X)^{-1} X^T y$$

It has a solution only when  $(X^T X)^{-1}$  is invertible (when its determinant is non-zero).

# ML (Example):

Let's test the normal equation by generating random data that follow linear pattern:

```
import numpy as np
X = 2 * np.random.rand(100, 1)
y = 4 + 3 * X + np.random.randn(100, 1)
```

Predictor

Target Variable


X

```
array([[0.10797752],
       [0.72385904],
       [0.42813739],
       [1.41056299],
       [1.28648117],
       [0.21776623],
       [1.03717871],
       [0.365245  ],
       [0.20016469],
       [0.20727274],
```

y

```
array([[ 3.5573093 ],
       [ 7.87380775],
       [ 7.2598704 ],
       [ 6.19588811],
       [ 9.13845766],
       [ 4.23094532],
       [ 8.61517587],
       [ 4.21443654],
       [ 6.3025794 ],
       [ 4.38441334],
```

# ML (Example):





# ML (Example):

```
X_b = np.c_[np.ones((100, 1)), X] # add x0 = 1 to each instance  
theta_best = np.linalg.inv(X_b.T.dot(X_b)).dot(X_b.T.dot(y))
```

Create matrix X

Predictor

X\_b

```
array([[1., 0.10797752],  
       [1., 0.72385904],  
       [1., 0.42813739],  
       [1., 1.41056299],  
       [1., 1.28648117],  
       [1., 0.21776623],  
       [1., 1.03717871],  
       [1., 0.365245  ],  
       [1., 0.20016469],  
       [1., 0.20727274],  
       ...])
```

theta\_best

```
array([[4.06669028],  
       [2.9236695 ]])
```

# ML (Example):

```
X_new = np.array([[0], [2]])  
X_new_b = np.c_[np.ones((2, 1)), X_new]  
y_predict = X_new_b.dot(theta_best)  
y_predict
```

```
array([[4.06669028],  
       [9.91402929]])
```

# ML (Exercise):

Calculate normal equation for the dataset of GDP per capita / Life satisfaction.



# ML (Gradient Descent):


➤ Problems with normal equations:

1. In many real cases  $(X^T X)^{-1}$  is not invertible,
2. Even if it is for big data sets the computational cost is  $O(n^3)$  or  $O(n^{2.4})$ .

➤ So, instead of calculating  $\theta$  from the normal equation, the learning algorithms use a technique to estimate this value which is called **Gradient Descent**.


1. Start with some initial parameters  $\theta$ ,
2. Tweaking the parameters ( $\theta$ ) iteratively, in a way that it reduces the cost function.

# ML (Gradient Descent-Example):




1. Start from a position  $(x, y)$ ,
2. Find the negative slope (to descend)
3. Take appropriate size step.


# ML (Gradient Descent-Example):




$$\begin{aligned}\theta_0 &= -1 \\ \theta_1 &= 1\end{aligned}$$



# ML (Gradient Descent-Example):




$$\begin{aligned}\theta_0 &= -0.5 \\ \theta_1 &= 1\end{aligned}$$






# ML (Gradient Descent-Example):



$$\theta_0 = 0.5$$
$$\theta_1 = 0.6$$



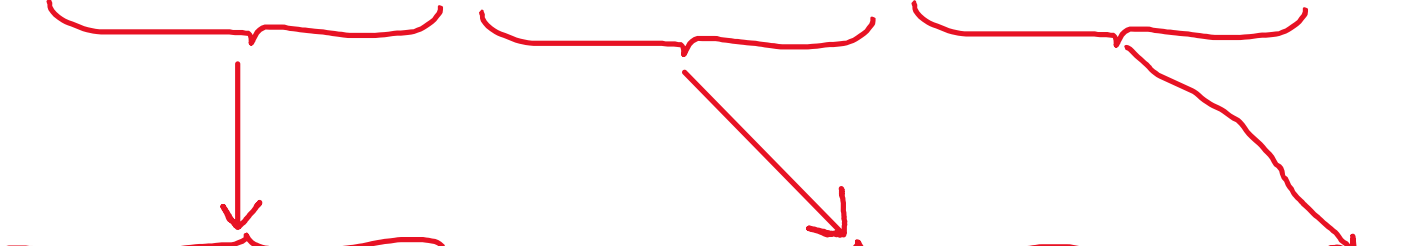


# ML (Gradient Descent-Example):

- Suppose we start with initial  $\theta_0$  and  $\theta_1$ .
- In which **direction** we should go to reduce the cost function?

$$MSE = (h_{\theta}(x^{(1)}) - y^{(1)})^2 + (h_{\theta}(x^{(2)}) - y^{(2)})^2 + (h_{\theta}(x^{(3)}) - y^{(3)})^2$$

$$\frac{d(MSE)}{d(\theta_0)} = \frac{d(MSE)}{d(h_{\theta})} \times \frac{d(h_{\theta})}{d(\theta_0)}$$


$$\frac{d(MSE)}{d(\theta_0)} = 2(h_{\theta}(x^{(1)}) - y^{(1)}) \times 1 + 2(h_{\theta}(x^{(2)}) - y^{(2)}) \times 1 + 2(h_{\theta}(x^{(3)}) - y^{(3)}) \times 1$$

$$\frac{d(MSE)}{d(\theta_1)} = \frac{d(MSE)}{d(h_{\theta})} \times \frac{d(h_{\theta})}{d(\theta_1)}$$

$$\frac{d(MSE)}{d(\theta_1)} = 2(h_{\theta}(x^{(1)}) - y^{(1)}) \times x^{(1)} + 2(h_{\theta}(x^{(2)}) - y^{(2)}) \times x^{(1)} + 2(h_{\theta}(x^{(3)}) - y^{(3)}) \times x^{(1)}$$

# ML (Gradient Descent-Example):

➤ Step sizes:  $\frac{d(MSE)}{d(\theta_0)}$   
Step size ( $\theta_0$ ) = slope  $\times$  *learning rate*

Step size ( $\theta_1$ ) = slope  $\times$  *learning rate*  
 $\frac{d(MSE)}{d(\theta_1)}$

Update Functions:

new  $\theta_0$  = previous  $\theta_0$  - step size  $\theta_0$

new  $\theta_1$  = previous  $\theta_1$  - step size  $\theta_1$

# ML (Gradient Descent-summary):

1. Random initialization of parameters ( $\theta_0, \theta_1, \theta_2, \dots$ ),
2. Calculate the slopes using gradient descent and **chain rule**,
3. Compute the steps using a pre-defined **learning rate**,
4. Update the parameters,

**Note:** In Machine learning, all the parameters that should be pre-defined in order to use the model are called **hyper parameters**. Hyper parameters are different from parameters. The parameters will be learned during training,...

# ML (Error Calculation):



Calculate the Root Mean Square Error for the predictions you made using linear regression for the dataframe GDP per capita/ life satisfaction(Use the functionalities of numpy).



# ML (Error Calculation):

Sub-module **metrics** and the function **mean\_squared\_error** of Sklearn.

```
from sklearn.metrics import mean_squared_error
mse = mean_squared_error(np.array(data['Life satisfaction']), predictions)
rmse = np.sqrt(mse)
```


# ML (Question):

- Is this error reliable for future predictions?
- Does it mean that our model will perform the best to predict?




# ML (Challenges of training):

## Overfitting



# ML (Challenges of training):

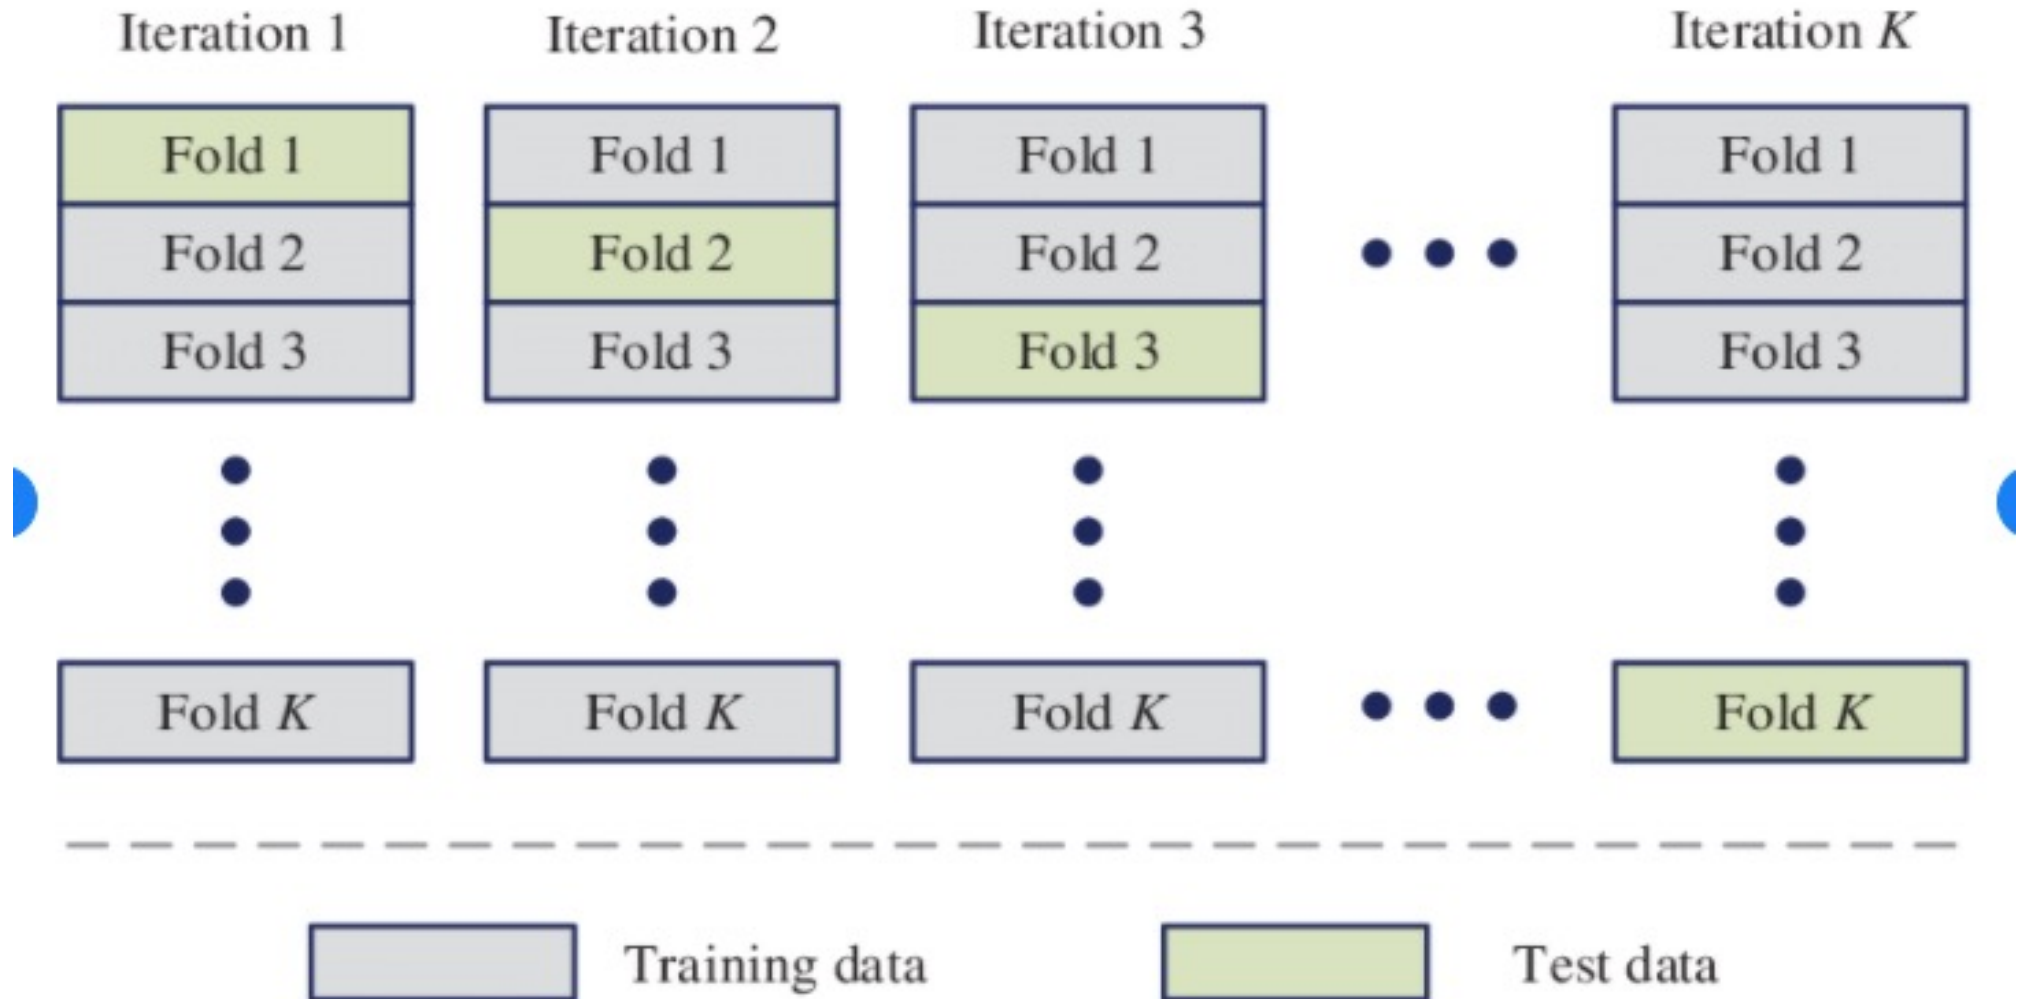
## Overfitting



What would be the prediction value for this point?



# ML (Concept of cross validation):



K-fold cross-validation method.

# ML (Concept of cross validation):

```
from sklearn.model_selection import cross_val_score
scores = cross_val_score(model, data[['GDP per capita']], data[['Life satisfaction']],
                          scoring="neg_mean_squared_error", cv=10)
rmse_scores = np.sqrt(-scores)
```

rmse\_scores

```
array([2.06550957, 0.98307636, 1.45811595, 2.0309329 , 3.211669  ,
       3.09781241, 1.78440725, 4.78091017, 2.29082548, 2.46068562])
```